

Week 3 - vapply and tapply

Sally Longmore

29/01/2020

In the last lesson, you learned about the two most fundamental members of R's *apply family of functions: `lapply()` and `sapply()`. Both take a list as input, apply a function to each element of the list, then combine and return the result. `lapply()` always returns a list, whereas `sapply` attempts to simplify the result.

In this lesson, you'll learn how to use **`vapply()`** and **`tapply()`**, each of which serves a very specific purpose within the Split-Apply-Combine methodology. For consistency, we'll use the same dataset we used in the **`lapply` and `sapply`** lesson.

The Flags dataset will be used again, which is stored in a variable called **`flags`**.

As you saw in the last lesson, the `unique()` function returns a vector of unique values contained in the object passed to it. Therefore, `sapply(flags, unique)` returns a list containing one vector of unique values for each column of the flags dataset.

```
sapply(flags, unique)
```

```
## $name
##   [1] Afghanistan      Albania            Algeria
##   [4] American-Samoa     Andorra           Angola
##   [7] Anguilla           Antigua-Barbuda   Argentina
##  [10] Argentine          Australia         Austria
##  [13] Bahamas           Bahrain           Bangladesh
##  [16] Barbados           Belgium           Belize
##  [19] Benin              Bermuda           Bhutan
##  [22] Bolivia            Botswana          Brazil
##  [25] British-Virgin-Isles Brunei            Bulgaria
##  [28] Burkina            Burma             Burundi
##  [31] Cameroon           Canada            Cape-Verde-Islands
##  [34] Cayman-Islands     Central-African-Republic Chad
##  [37] Chile              China             Colombia
##  [40] Comorro-Islands    Congo             Cook-Islands
##  [43] Costa-Rica         Cuba             Cyprus
##  [46] Czechoslovakia     Denmark           Djibouti
##  [49] Dominica           Dominican-Republic Ecuador
##  [52] Egypt             El-Salvador       Equatorial-Guinea
##  [55] Ethiopia           Faeroes           Falklands-Malvinas
##  [58] Fiji              Finland           France
##  [61] French-Guiana      French-Polynesia  Gabon
##  [64] Gambia             Germany-DDR       Germany-FRG
##  [67] Ghana             Gibraltar         Greece
##  [70] Greenland          Grenada           Guam
##  [73] Guatemala          Guinea            Guinea-Bissau
##  [76] Guyana             Haiti             Honduras
##  [79] Hong-Kong          Hungary           Iceland
```

## [82]	India	Indonesia	Iran
## [85]	Iraq	Ireland	Israel
## [88]	Italy	Ivory-Coast	Jamaica
## [91]	Japan	Jordan	Kampuchea
## [94]	Kenya	Kiribati	Kuwait
## [97]	Laos	Lebanon	Lesotho
## [100]	Liberia	Libya	Liechtenstein
## [103]	Luxembourg	Malagasy	Malawi
## [106]	Malaysia	Maldives-Islands	Mali
## [109]	Malta	Marianas	Mauritania
## [112]	Mauritius	Mexico	Micronesia
## [115]	Monaco	Mongolia	Montserrat
## [118]	Morocco	Mozambique	Nauru
## [121]	Nepal	Netherlands	Netherlands-Antilles
## [124]	New-Zealand	Nicaragua	Niger
## [127]	Nigeria	Niue	North-Korea
## [130]	North-Yemen	Norway	Oman
## [133]	Pakistan	Panama	Papua-New-Guinea
## [136]	Parguay	Peru	Philippines
## [139]	Poland	Portugal	Puerto-Rico
## [142]	Qatar	Romania	Rwanda
## [145]	San-Marino	Sao-Tome	Saudi-Arabia
## [148]	Senegal	Seychelles	Sierra-Leone
## [151]	Singapore	Soloman-Islands	Somalia
## [154]	South-Africa	South-Korea	South-Yemen
## [157]	Spain	Sri-Lanka	St-Helena
## [160]	St-Kitts-Nevis	St-Lucia	St-Vincent
## [163]	Sudan	Surinam	Swaziland
## [166]	Sweden	Switzerland	Syria
## [169]	Taiwan	Tanzania	Thailand
## [172]	Togo	Tonga	Trinidad-Tobago
## [175]	Tunisia	Turkey	Turks-Cocos-Islands
## [178]	Tuvalu	UAE	Uganda
## [181]	UK	Uruguay	US-Virgin-Isles
## [184]	USA	USSR	Vanuatu
## [187]	Vatican-City	Venezuela	Vietnam
## [190]	Western-Samoa	Yugoslavia	Zaire
## [193]	Zambia	Zimbabwe	

194 Levels: Afghanistan Albania Algeria American-Samoa Andorra ... Zimbabwe

##

\$landmass

[1] 5 3 4 6 1 2

##

\$zone

[1] 1 3 2 4

##

\$area

## [1]	648	29	2388	0	1247	2777	7690	84	19	1	143	31
## [13]	23	113	47	1099	600	8512	6	111	274	678	28	474
## [25]	9976	4	623	1284	757	9561	1139	2	342	51	115	9
## [37]	128	43	22	49	284	1001	21	1222	12	18	337	547
## [49]	91	268	10	108	249	239	132	2176	109	246	36	215
## [61]	112	93	103	3268	1904	1648	435	70	301	323	11	372
## [73]	98	181	583	236	30	1760	3	587	118	333	1240	1031

```

## [85] 1973 1566 447 783 140 41 1267 925 121 195 324 212
## [97] 804 76 463 407 1285 300 313 92 237 26 2150 196
## [109] 72 637 1221 99 288 505 66 2506 63 17 450 185
## [121] 945 514 57 5 164 781 245 178 9363 22402 15 912
## [133] 256 905 753 391
##
## $population
## [1] 16 3 20 0 7 28 15 8 90 10 1 6 119 9 35
## [16] 4 24 2 11 1008 5 47 31 54 17 61 14 684 157 39
## [31] 57 118 13 77 12 56 18 84 48 36 22 29 38 49 45
## [46] 231 274 60
##
## $language
## [1] 10 6 8 1 2 4 3 5 7 9
##
## $religion
## [1] 2 6 1 0 5 3 4 7
##
## $bars
## [1] 0 2 3 1 5
##
## $stripes
## [1] 3 0 2 1 5 9 11 14 4 6 13 7
##
## $colours
## [1] 5 3 2 8 6 4 7 1
##
## $red
## [1] 1 0
##
## $green
## [1] 1 0
##
## $blue
## [1] 0 1
##
## $gold
## [1] 1 0
##
## $white
## [1] 1 0
##
## $black
## [1] 1 0
##
## $orange
## [1] 0 1
##
## $mainhue
## [1] green red blue gold white orange black brown
## Levels: black blue brown gold green orange red white
##
## $circles
## [1] 0 1 4 2

```

```
##
## $crosses
## [1] 0 1 2
##
## $saltires
## [1] 0 1
##
## $quarters
## [1] 0 1 4
##
## $sunstars
## [1] 1 0 6 22 14 3 4 5 15 10 7 2 9 50
##
## $crescent
## [1] 0 1
##
## $triangle
## [1] 0 1
##
## $icon
## [1] 1 0
##
## $animate
## [1] 0 1
##
## $text
## [1] 0 1
##
## $topleft
## [1] black red green blue white orange gold
## Levels: black blue gold green orange red white
##
## $botright
## [1] green red white black blue gold orange brown
## Levels: black blue brown gold green orange red white
```

What if you had forgotten how `unique()` works and mistakenly thought it returns the **number** of unique values contained in the object passed to it? Then you might have incorrectly expected `sapply(flags, unique)` to return a numeric vector, since each element of the list returned would contain a single number and `sapply()` could then simplify the result to a vector.

When working interactively (at the prompt), this is not much of a problem, since you see the result immediately and will quickly recognize your mistake. However, when working non-interactively (e.g. writing your own functions), a misunderstanding may go undetected and cause incorrect results later on. Therefore, you may wish to be more careful and that's where `vapply()` is useful.

Whereas `sapply` tries to *guess* the correct format of the result, `vapply()` allows you to specify it explicitly. If the result doesn't match the format you specify, `vapply()` will throw an error, causing the operation to stop. This can prevent significant problems in your code that might be caused by getting an unexpected return value from `sapply()`.

Try `vapply(flags, unique, numeric(1))`, which says that you expect each element of the result to be a numeric vector of length 1. Since this is NOT actually the case, **YOU WILL GET AN ERROR**. Once you get the error, type `ok()` to continue to the next question.

```
vapply(flags, unique, numeric(1))
```

Recall from the previous lesson that `sapply(flags, class)` will return a character vector containing the class of each column in the dataset. Try that again now to see the result.

```
sapply(flags, class)
```

```
##      name  landmass      zone      area population  language  religion
## "factor" "integer" "integer" "integer" "integer" "integer" "integer"
##      bars  stripes  colours      red      green      blue      gold
## "integer" "integer" "integer" "integer" "integer" "integer" "integer"
##      white  black  orange  mainhue  circles  crosses  saltires
## "integer" "integer" "integer" "factor" "integer" "integer" "integer"
## quarters  sunstars  crescent  triangle  icon  animate  text
## "integer" "integer" "integer" "integer" "integer" "integer" "integer"
## topleft  botright
## "factor" "factor"
```

If we wish to be explicit about the format of the result we expect, we can use `vapply(flags, class, character(1))`. The **character(1)** argument tells R that we expect the class function to return a character vector of length 1 when applied to **EACH** column of the flags dataset.

```
vapply(flags, class, character(1))
```

```
##      name  landmass      zone      area population  language  religion
## "factor" "integer" "integer" "integer" "integer" "integer" "integer"
##      bars  stripes  colours      red      green      blue      gold
## "integer" "integer" "integer" "integer" "integer" "integer" "integer"
##      white  black  orange  mainhue  circles  crosses  saltires
## "integer" "integer" "integer" "factor" "integer" "integer" "integer"
## quarters  sunstars  crescent  triangle  icon  animate  text
## "integer" "integer" "integer" "integer" "integer" "integer" "integer"
## topleft  botright
## "factor" "factor"
```

Note that since our expectation was correct (i.e. `character(1)`), the `vapply()` result is identical to the `sapply()` result – a character vector of column classes.

You may think of `vapply()` as being **safer** than `sapply()`, since it requires you to specify the format of the output in advance, instead of just allowing R to **guess** what you wanted. In addition, `vapply()` may perform faster than `sapply` for large datasets. However, when doing data analysis interactively, `sapply()` saves you some typing and will often be good enough.

As a data analyst, you'll often wish to split your data up into groups based on the value of some variable, then apply a function to the members of each group. The next function we will look at **tapply()** does exactly that.

The **landmass** variable in our dataset takes integer values between 1 and 6, each of which represents a different part of the world. Use `table(flags$landmass)` to see how many flags/countries fall into each group.

```
table(flags$landmass)
```

```
##
##  1  2  3  4  5  6
## 31 17 35 52 39 20
```

The **animate** variable in our dataset takes the value 1 if a countries flag contains an animate image (e.g. an eagle, a tree, a human hand) and 0 otherwise. Use `table(flags$animate)` to see how many flags contain an animate imate.

```
table(flags$animate)
```

```
##
##    0    1
## 155   39
```

This tells us that 39 flags contain an animate object (animate = 1) and 155 do not (animate = 0).

If you take the arithmetic mean of a bunch of 0s and 1s, you get the proportion of 1s. Use the `tapply(flags$animate, flags$landmass, mean)` to apply the mean function to the **animate** variable separately for each of the six landmass groups, this giving is the proportion of flags containing an animate image WITHIN each landmass group.

```
tapply(flags$animate, flags$landmass, mean)
```

```
##           1           2           3           4           5           6
## 0.4193548 0.1764706 0.1142857 0.1346154 0.1538462 0.3000000
```

The first landmass group (landmass = 1) corresponds to North America and contains the highest proportion of flags with animate image (0.4194).

Similarly, we can look at a summary of population values (in round millions) for countries with and without the colour red on their flag with `tapply(flags$population, flags$red, summary)`.

```
tapply(flags$population, flags$red, summary)
```

```
## $`0`
##   Min. 1st Qu.  Median    Mean 3rd Qu.    Max.
##   0.00   0.00    3.00   27.63   9.00   684.00
##
## $`1`
##   Min. 1st Qu.  Median    Mean 3rd Qu.    Max.
##   0.0    0.0    4.0    22.1   15.0  1008.0
```

What is the median population (in millions) for countries **without** the colour red on their flag?

3.0

Lastly, use the same approach to look at a summary of population values for each of the six landmasses.

```
tapply(flags$population, flags$landmass, summary)
```

```
## $`1`
##   Min. 1st Qu.  Median    Mean 3rd Qu.    Max.
##   0.00   0.00   0.00   12.29   4.50   231.00
##
## $`2`
##   Min. 1st Qu.  Median    Mean 3rd Qu.    Max.
##   0.00   1.00   6.00   15.71   15.00   119.00
##
## $`3`
##   Min. 1st Qu.  Median    Mean 3rd Qu.    Max.
##   0.00   0.00   8.00   13.86   16.00   61.00
##
## $`4`
##   Min. 1st Qu.  Median    Mean 3rd Qu.    Max.
##   0.000   1.000   5.000   8.788   9.750   56.000
##
## $`5`
```

```
##      Min. 1st Qu.  Median      Mean 3rd Qu.      Max.
##      0.00    2.00   10.00   69.18   39.00 1008.00
##
## $`6`
##      Min. 1st Qu.  Median      Mean 3rd Qu.      Max.
##      0.00    0.00    0.00   11.30    1.25   157.00
```

What is the maximum population (in millions) for the fourth landmass group (Africa)?

56.0