

Week 4 - Looking at Data

Sally Longmore

30/01/2020

Whenever you are working with a new dataset, the first thing you should do is look at it and ask yourself some questions. - What is the format of the data? - What are the dimensions? - What are the variable names? - How are the variables stored? - Are there missing data? - Are there any flaws in the data?

This lesson will teach you how to answer these questions and more using R's built-in functions. We will be using a dataset constructed from the United States Department of Agriculture's PLANTS database.

(http://plants.usda.gov/adv_search.html)

The data is stored in a variable called plants.

```
plants <- readRDS("plants.rds")
```

Type ls() to list the variables in your workspace, among which should be plants.

```
ls()
```

```
## [1] "plants"
```

Let's begin by checking the class of the plants variable with class(plants). This will give us a clue as to the overall structure of the data.

```
class(plants)
```

```
## [1] "data.frame"
```

It's very common for data to be stored in a data frame. It is the default class for data read into R using functions like read.csv() and read.table(), which you will learn about in another lesson.

Since the dataset is stored in a data frame, we know it is rectangular. In other words, it has two dimensions (rows and columns) and fits neatly into a table or spreadsheet. Use dim(plants) to see exactly how many rows and columns we are dealing with.

```
dim(plants)
```

```
## [1] 5166 10
```

The first number you see (5166) is the number of rows (observations) and the second number (10) is the number of columns (variables).

You can also use nrow(plants) to see only the number of rows.

```
nrow(plants)
```

```
## [1] 5166
```

And you can run ncol(plants) to see the number of columns.

```
ncol(plants)
```

```
## [1] 10
```

If you are curious about how much space the dataset is occupying in memory, you can use `object.size(plants)`.

```
object.size(plants)
```

```
## 686080 bytes
```

Now that we have a sense of the shape and size of the dataset, let's get a feel for what is inside. we can use `names(plants)` to return a character vector of column (i.e. variable) names.

```
names(plants)
```

```
## [1] "Scientific_Name"      "Duration"             "Active_Growth_Period"
## [4] "Foliage_Color"       "pH_Min"               "pH_Max"
## [7] "Precip_Min"          "Precip_Max"           "Shade_Tolerance"
## [10] "Temp_Min_F"
```

We have applied fairly descriptive variable names to this dataset, but that won't always be the case. A logical next step is to peak at the actual data. However, our dataset contains over 5,000 observations (rows), so it is impractical to view the whole dataset at once.

The `head()` function allows you to preview the top of the dataset.

```
head(plants)
```

```
##           Scientific_Name      Duration Active_Growth_Period
## 1           Abielmoschus          <NA>          <NA>
## 2    Abielmoschus esculentus Annual, Perennial          <NA>
## 3                Abies          <NA>          <NA>
## 4    Abies balsamea      Perennial    Spring and Summer
## 5 Abies balsamea var. balsamea      Perennial          <NA>
## 6           Abutilon          <NA>          <NA>
##  Foliage_Color pH_Min pH_Max Precip_Min Precip_Max Shade_Tolerance Temp_Min_F
## 1          <NA>    NA    NA         NA         NA          <NA>         NA
## 2          <NA>    NA    NA         NA         NA          <NA>         NA
## 3          <NA>    NA    NA         NA         NA          <NA>         NA
## 4         Green     4     6         13         60        Tolerant        -43
## 5          <NA>    NA    NA         NA         NA          <NA>         NA
## 6          <NA>    NA    NA         NA         NA          <NA>         NA
```

Take a minute to look through and understand the output above. Each row is labelled with the observation number and each column with the variable name. Your screen is probably not wide enough to view all 10 columns side-by-side, in which case R displays as many columns as it can on each line before continuing to the next.

By default, `head()` shows you the first six rows of data. You can alter this behaviour by passing as a second argument the number of rows you would like to view. Use `head()` to preview the first 10 rows of `plants`.

```
head(plants, 10)
```

```
##           Scientific_Name      Duration Active_Growth_Period
## 1           Abielmoschus          <NA>          <NA>
## 2    Abielmoschus esculentus Annual, Perennial          <NA>
## 3                Abies          <NA>          <NA>
## 4    Abies balsamea      Perennial    Spring and Summer
## 5 Abies balsamea var. balsamea      Perennial          <NA>
## 6           Abutilon          <NA>          <NA>
## 7    Abutilon theophrasti      Annual          <NA>
## 8                Acacia          <NA>          <NA>
## 9    Acacia constricta      Perennial    Spring and Summer
```

```
## 10 Acacia constricta var. constricta      Perennial      <NA>
##      Foliage_Color pH_Min pH_Max Precip_Min Precip_Max Shade_Tolerance Temp_Min_F
## 1      <NA>      NA      NA      NA      NA      <NA>      NA
## 2      <NA>      NA      NA      NA      NA      <NA>      NA
## 3      <NA>      NA      NA      NA      NA      <NA>      NA
## 4      Green      4      6.0      13      60      Tolerant      -43
## 5      <NA>      NA      NA      NA      NA      <NA>      NA
## 6      <NA>      NA      NA      NA      NA      <NA>      NA
## 7      <NA>      NA      NA      NA      NA      <NA>      NA
## 8      <NA>      NA      NA      NA      NA      <NA>      NA
## 9      Green      7      8.5      4      20      Intolerant      -13
## 10     <NA>      NA      NA      NA      NA      <NA>      NA
```

The same applies for using `tail()` to preview the end of the dataset. Use `tail()` to view the last 15 rows.

```
tail(plants, 15)
```

```
##      Scientific_Name Duration Active_Growth_Period
## 5152      Zizania      <NA>      <NA>
## 5153      Zizania aquatica Annual      Spring
## 5154      Zizania aquatica var. aquatica Annual      <NA>
## 5155      Zizania palustris Annual      <NA>
## 5156      Zizania palustris var. palustris Annual      <NA>
## 5157      Zizaniopsis      <NA>      <NA>
## 5158      Zizaniopsis miliacea Perennial      Spring and Summer
## 5159      Zizia      <NA>      <NA>
## 5160      Zizia aptera Perennial      <NA>
## 5161      Zizia aurea Perennial      <NA>
## 5162      Zizia trifoliata Perennial      <NA>
## 5163      Zostera      <NA>      <NA>
## 5164      Zostera marina Perennial      <NA>
## 5165      Zoysia      <NA>      <NA>
## 5166      Zoysia japonica Perennial      <NA>
##      Foliage_Color pH_Min pH_Max Precip_Min Precip_Max Shade_Tolerance
## 5152     <NA>      NA      NA      NA      NA      <NA>
## 5153     Green      6.4      7.4      30      50      Intolerant
## 5154     <NA>      NA      NA      NA      NA      <NA>
## 5155     <NA>      NA      NA      NA      NA      <NA>
## 5156     <NA>      NA      NA      NA      NA      <NA>
## 5157     <NA>      NA      NA      NA      NA      <NA>
## 5158     Green      4.3      9.0      35      70      Intolerant
## 5159     <NA>      NA      NA      NA      NA      <NA>
## 5160     <NA>      NA      NA      NA      NA      <NA>
## 5161     <NA>      NA      NA      NA      NA      <NA>
## 5162     <NA>      NA      NA      NA      NA      <NA>
## 5163     <NA>      NA      NA      NA      NA      <NA>
## 5164     <NA>      NA      NA      NA      NA      <NA>
## 5165     <NA>      NA      NA      NA      NA      <NA>
## 5166     <NA>      NA      NA      NA      NA      <NA>
##      Temp_Min_F
## 5152      NA
## 5153      32
## 5154      NA
## 5155      NA
## 5156      NA
```

```
## 5157      NA
## 5158      12
## 5159      NA
## 5160      NA
## 5161      NA
## 5162      NA
## 5163      NA
## 5164      NA
## 5165      NA
## 5166      NA
```

After previewing the top and bottom of the data, you probably noticed lots of NAs, which are R's placeholders for missing values. Use `summary(plants)` to get a better feel for how each variable is distributed and how much of the dataset is missing.

```
summary(plants)
```

```
##           Scientific_Name           Duration
## Abelmoschus      : 1 Perennial      :3031
## Abelmoschus esculentus : 1 Annual      : 682
## Abies             : 1 Annual, Perennial: 179
## Abies balsamea     : 1 Annual, Biennial : 95
## Abies balsamea var. balsamea: 1 Biennial : 57
## Abutilon           : 1 (Other)        : 92
## (Other)            :5160 NA's         :1030
##           Active_Growth_Period      Foliage_Color      pH_Min
## Spring and Summer : 447 Dark Green : 82 Min. :3.000
## Spring             : 144 Gray-Green : 25 1st Qu.:4.500
## Spring, Summer, Fall: 95 Green      : 692 Median :5.000
## Summer             : 92 Red         : 4 Mean :4.997
## Summer and Fall    : 24 White-Gray : 9 3rd Qu.:5.500
## (Other)            : 30 Yellow-Green: 20 Max. :7.000
## NA's               :4334 NA's       :4334 NA's :4327
##           pH_Max      Precip_Min      Precip_Max      Shade_Tolerance
## Min. : 5.100 Min. : 4.00 Min. : 16.00 Intermediate: 242
## 1st Qu.: 7.000 1st Qu.:16.75 1st Qu.: 55.00 Intolerant : 349
## Median : 7.300 Median :28.00 Median : 60.00 Tolerant : 246
## Mean : 7.344 Mean :25.57 Mean : 58.73 NA's :4329
## 3rd Qu.: 7.800 3rd Qu.:32.00 3rd Qu.: 60.00
## Max. :10.000 Max. :60.00 Max. :200.00
## NA's :4327 NA's :4338 NA's :4338
##           Temp_Min_F
## Min. : -79.00
## 1st Qu.: -38.00
## Median : -33.00
## Mean : -22.53
## 3rd Qu.: -18.00
## Max. : 52.00
## NA's :4328
```

`summary()` provides different output for each variable, depending on its class. For numeric data such as **Precip_Min**, `summary()` displays the minimum, 1st quartile, median, mean, 3rd quartile and maximum. These values help us understand how the data are distributed.

For categorical variables (called 'factor' variables in R), `summary()` displays the number of times each value (or 'level') occurs in the data. For example, each value of **Scientific_Name** only appears once, since it is

unique to specific plants. In contrast, the summary for `Duration` (also a factor variable) tells us that our dataset contains 3031 Perennial plants, 682 Annual plants etc.

You can see that R truncated the summary for `Active_Growth_Period` by including a catch-all category called **Other**. Since it is a categorical/factor variable, we can see how many times each value actually occurs in the data with `table(plants$Active_Growth_Period)`.

```
table(plants$Active_Growth_Period)
```

```
##
## Fall, Winter and Spring      Spring      Spring and Fall
##           15              144             10
##      Spring and Summer  Spring, Summer, Fall      Summer
##           447             95             92
##      Summer and Fall      Year Round
##           24              5
```

Each of the functions we have introduced so far has its place in helping you to better understand the structure of your data. However, we have left the best for last.

Perhaps the most useful and concise function for understanding the **structure** of your data is the `str()` function.

```
str(plants)
```

```
## 'data.frame':   5166 obs. of  10 variables:
## $ Scientific_Name      : Factor w/ 5166 levels "Abelmoschus",...: 1 2 3 4 5 6 7 8 9 10 ...
## $ Duration             : Factor w/ 8 levels "Annual","Annual, Biennial",...: NA 4 NA 7 7 NA 1 NA 7 7
## $ Active_Growth_Period: Factor w/ 8 levels "Fall, Winter and Spring",...: NA NA NA 4 NA NA NA NA 4 NA
## $ Foliage_Color        : Factor w/ 6 levels "Dark Green","Gray-Green",...: NA NA NA 3 NA NA NA NA 3 NA
## $ pH_Min              : num  NA NA NA 4 NA NA NA NA 7 NA ...
## $ pH_Max              : num  NA NA NA 6 NA NA NA NA 8.5 NA ...
## $ Precip_Min          : int   NA NA NA 13 NA NA NA NA 4 NA ...
## $ Precip_Max          : int   NA NA NA 60 NA NA NA NA 20 NA ...
## $ Shade_Tolerance     : Factor w/ 3 levels "Intermediate",...: NA NA NA 3 NA NA NA NA 2 NA ...
## $ Temp_Min_F          : int   NA NA NA -43 NA NA NA NA -13 NA ...
```

The beauty of `str()` is that it combines many of the features of the other functions you have already seen, all in a concise and readable format. At the very top, it tells us that the class of `plants` is **data.frame** and that it has 5166 observations and 10 variables. It then gives us the name and class of each variable, as well as a preview of its contents.

`str()` is actually a very general function that you can use on most objects in R. Any time you want to understand the structure of something (a dataset, function, etc.), `str()` is a good place to start.